



AUSTIN DEEP LEARNING MEETUP

*John Ramey
@ramhiser
Principal Data Scientist, Novi Labs*

Welcome!

66

In deep learning, the algorithms we use now are versions of the algorithms we were developing in the 1980s, the 1990s. People were very optimistic about them, but it turns out they didn't work too well. Now we know the reason is they didn't work too well is that we didn't have powerful enough computers, we didn't have enough data sets to train them.

-Geoffrey Hinton

YISSSSSSS!!!

Austin Deep Learning

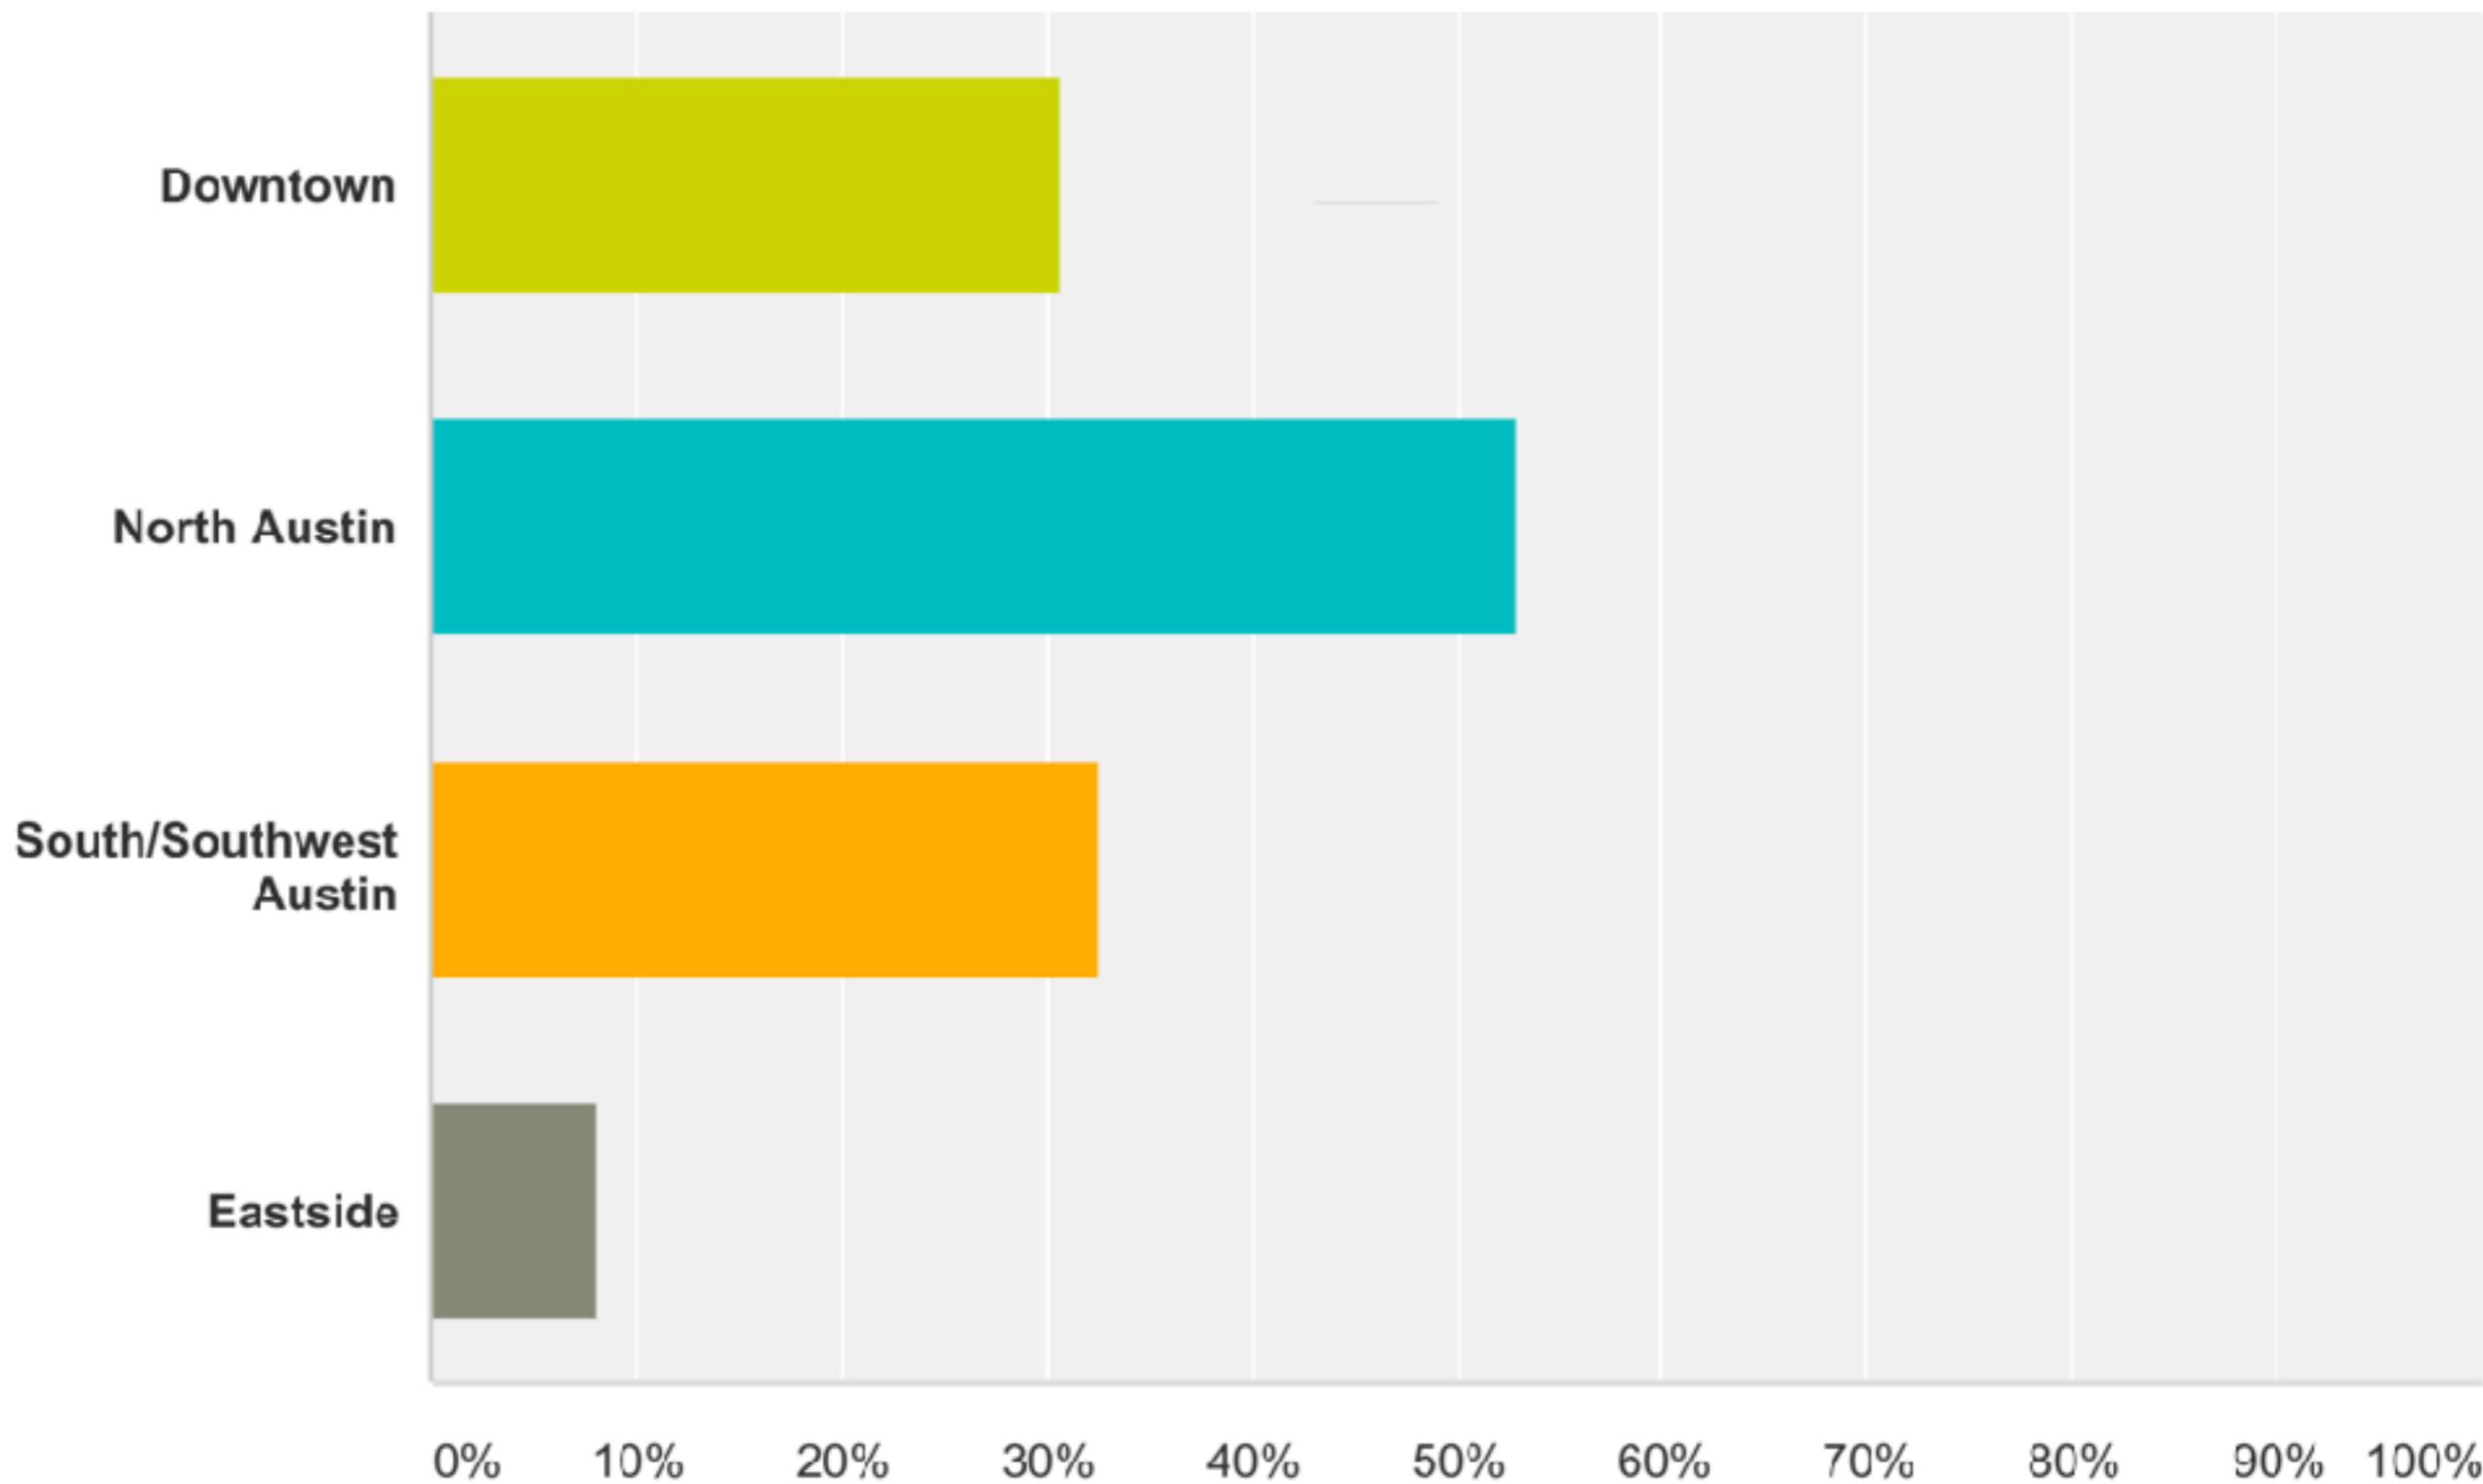
17 new members joined your Meetup Group!

You now have **506 members!** Nice work!

Survey Results

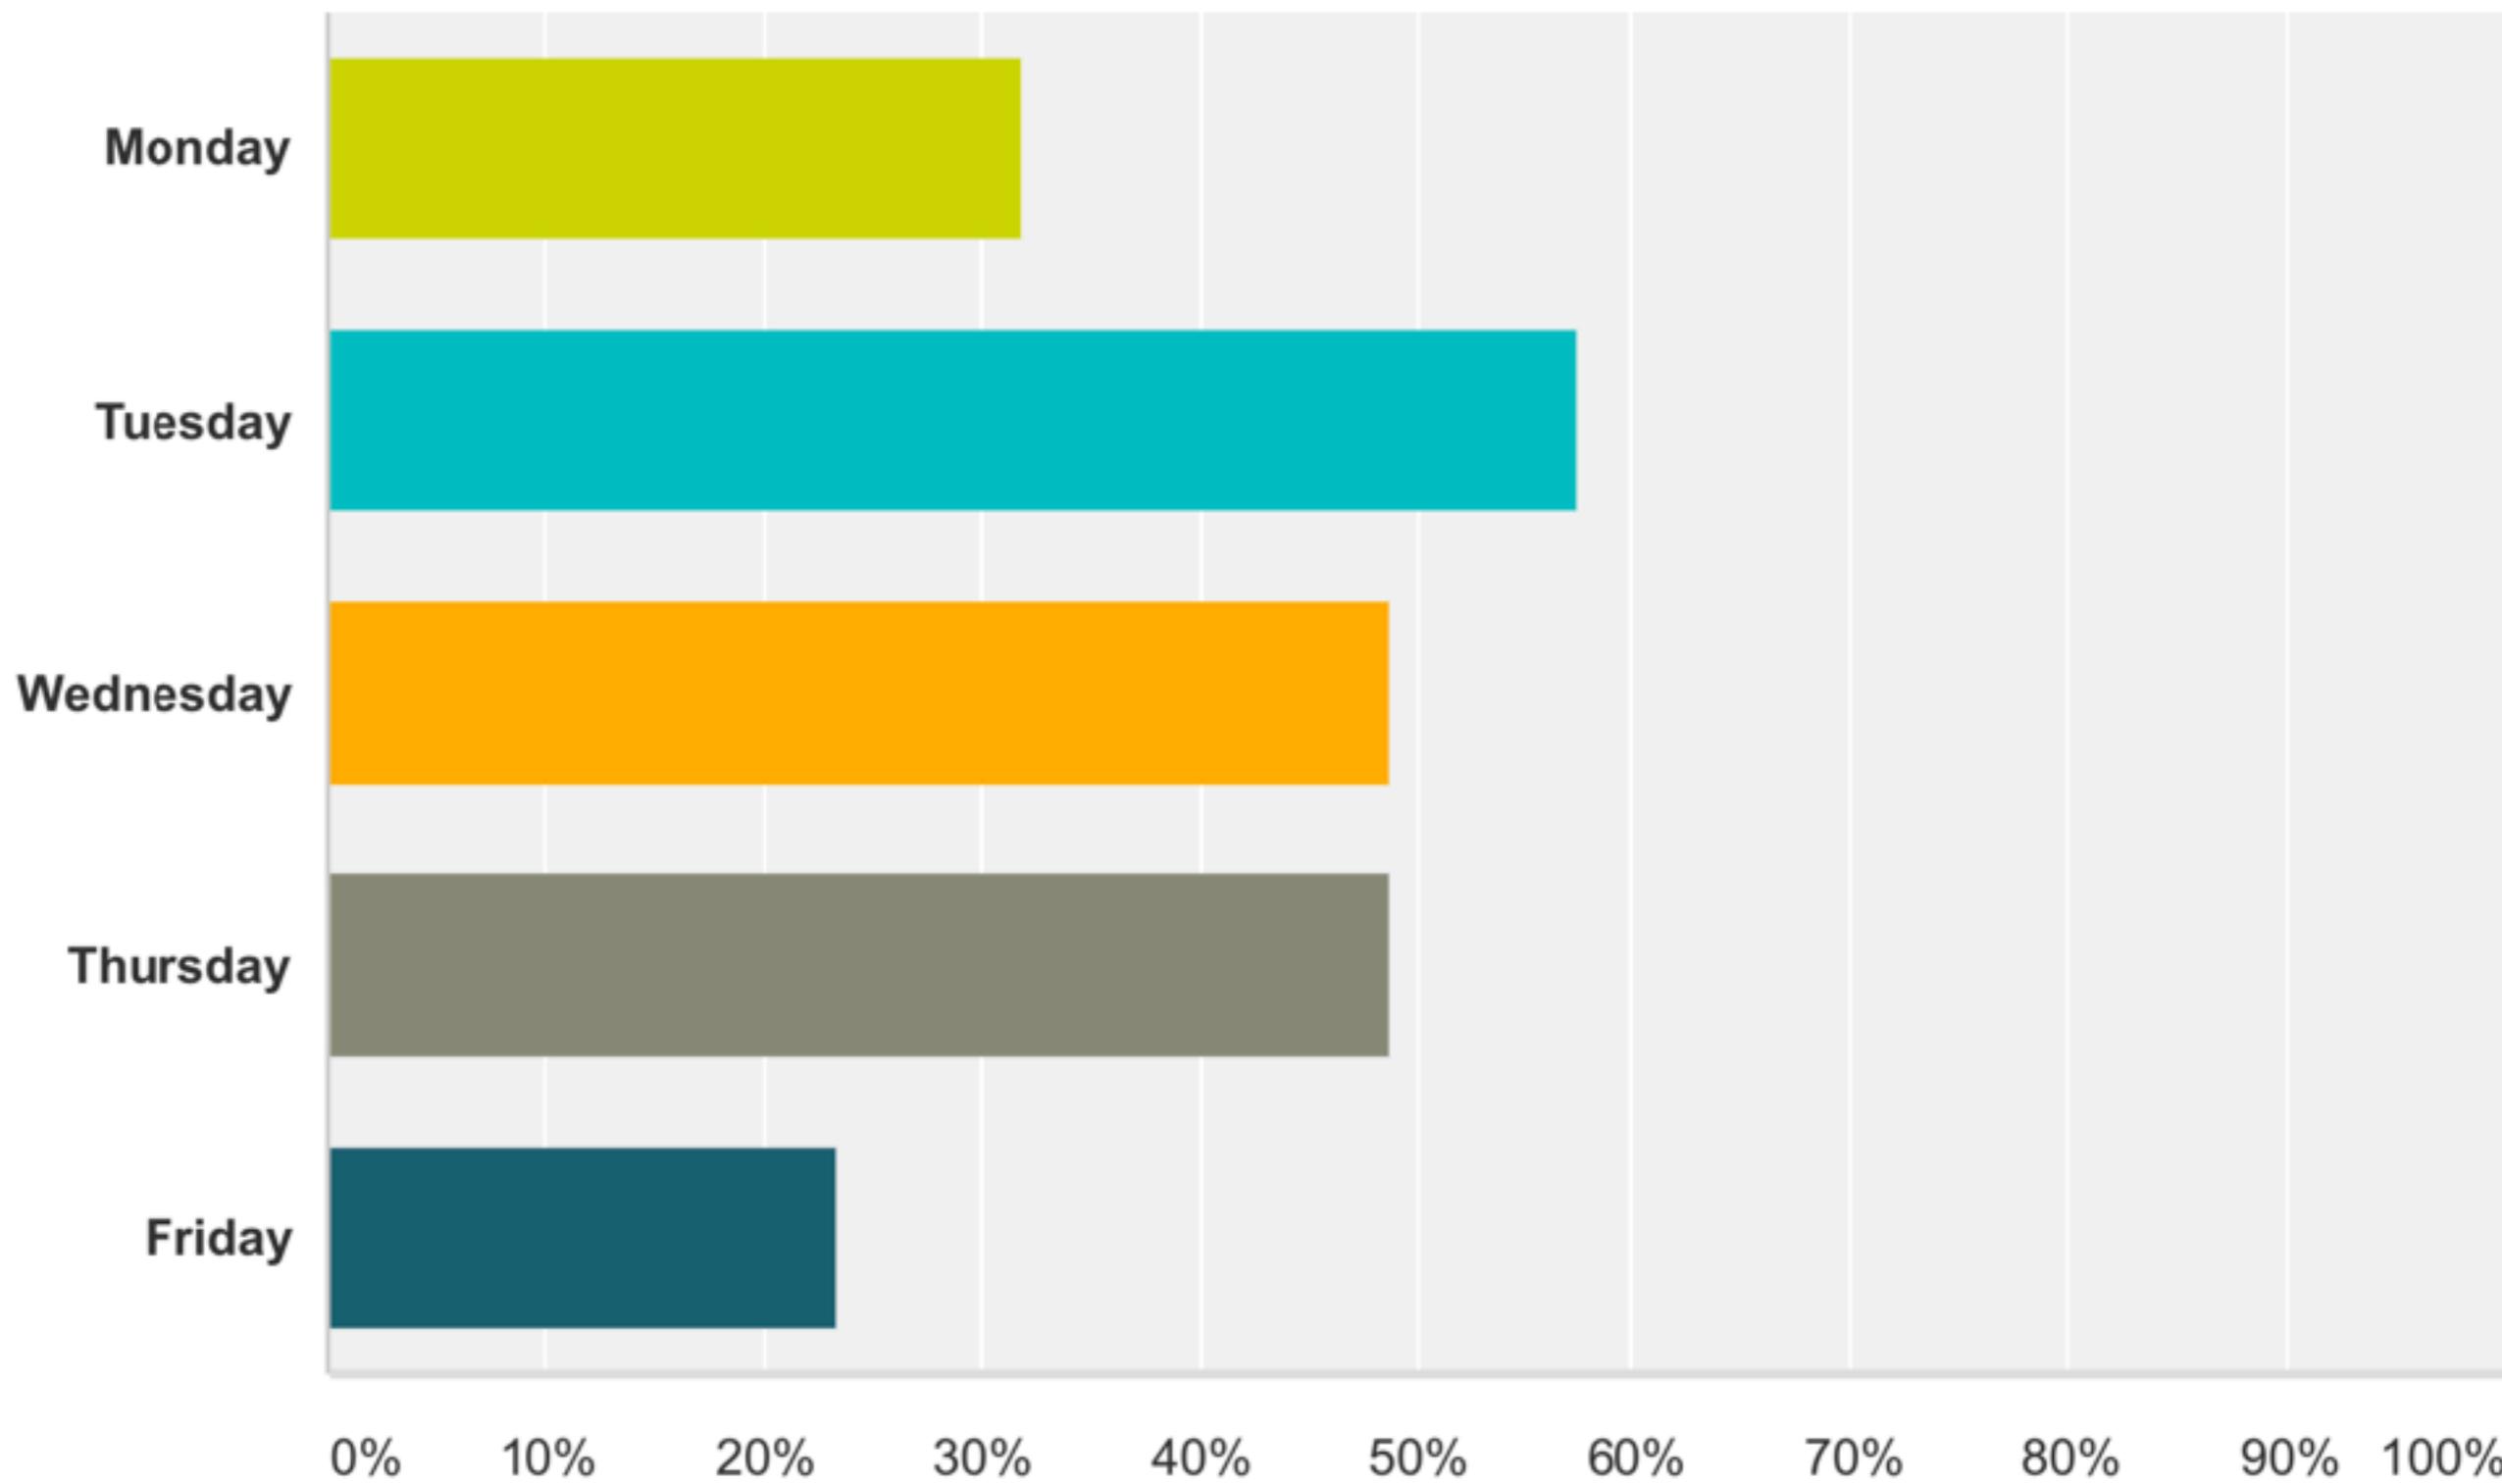
Which part of Austin would you prefer to have our meetups?

Answered: 49 Skipped: 0



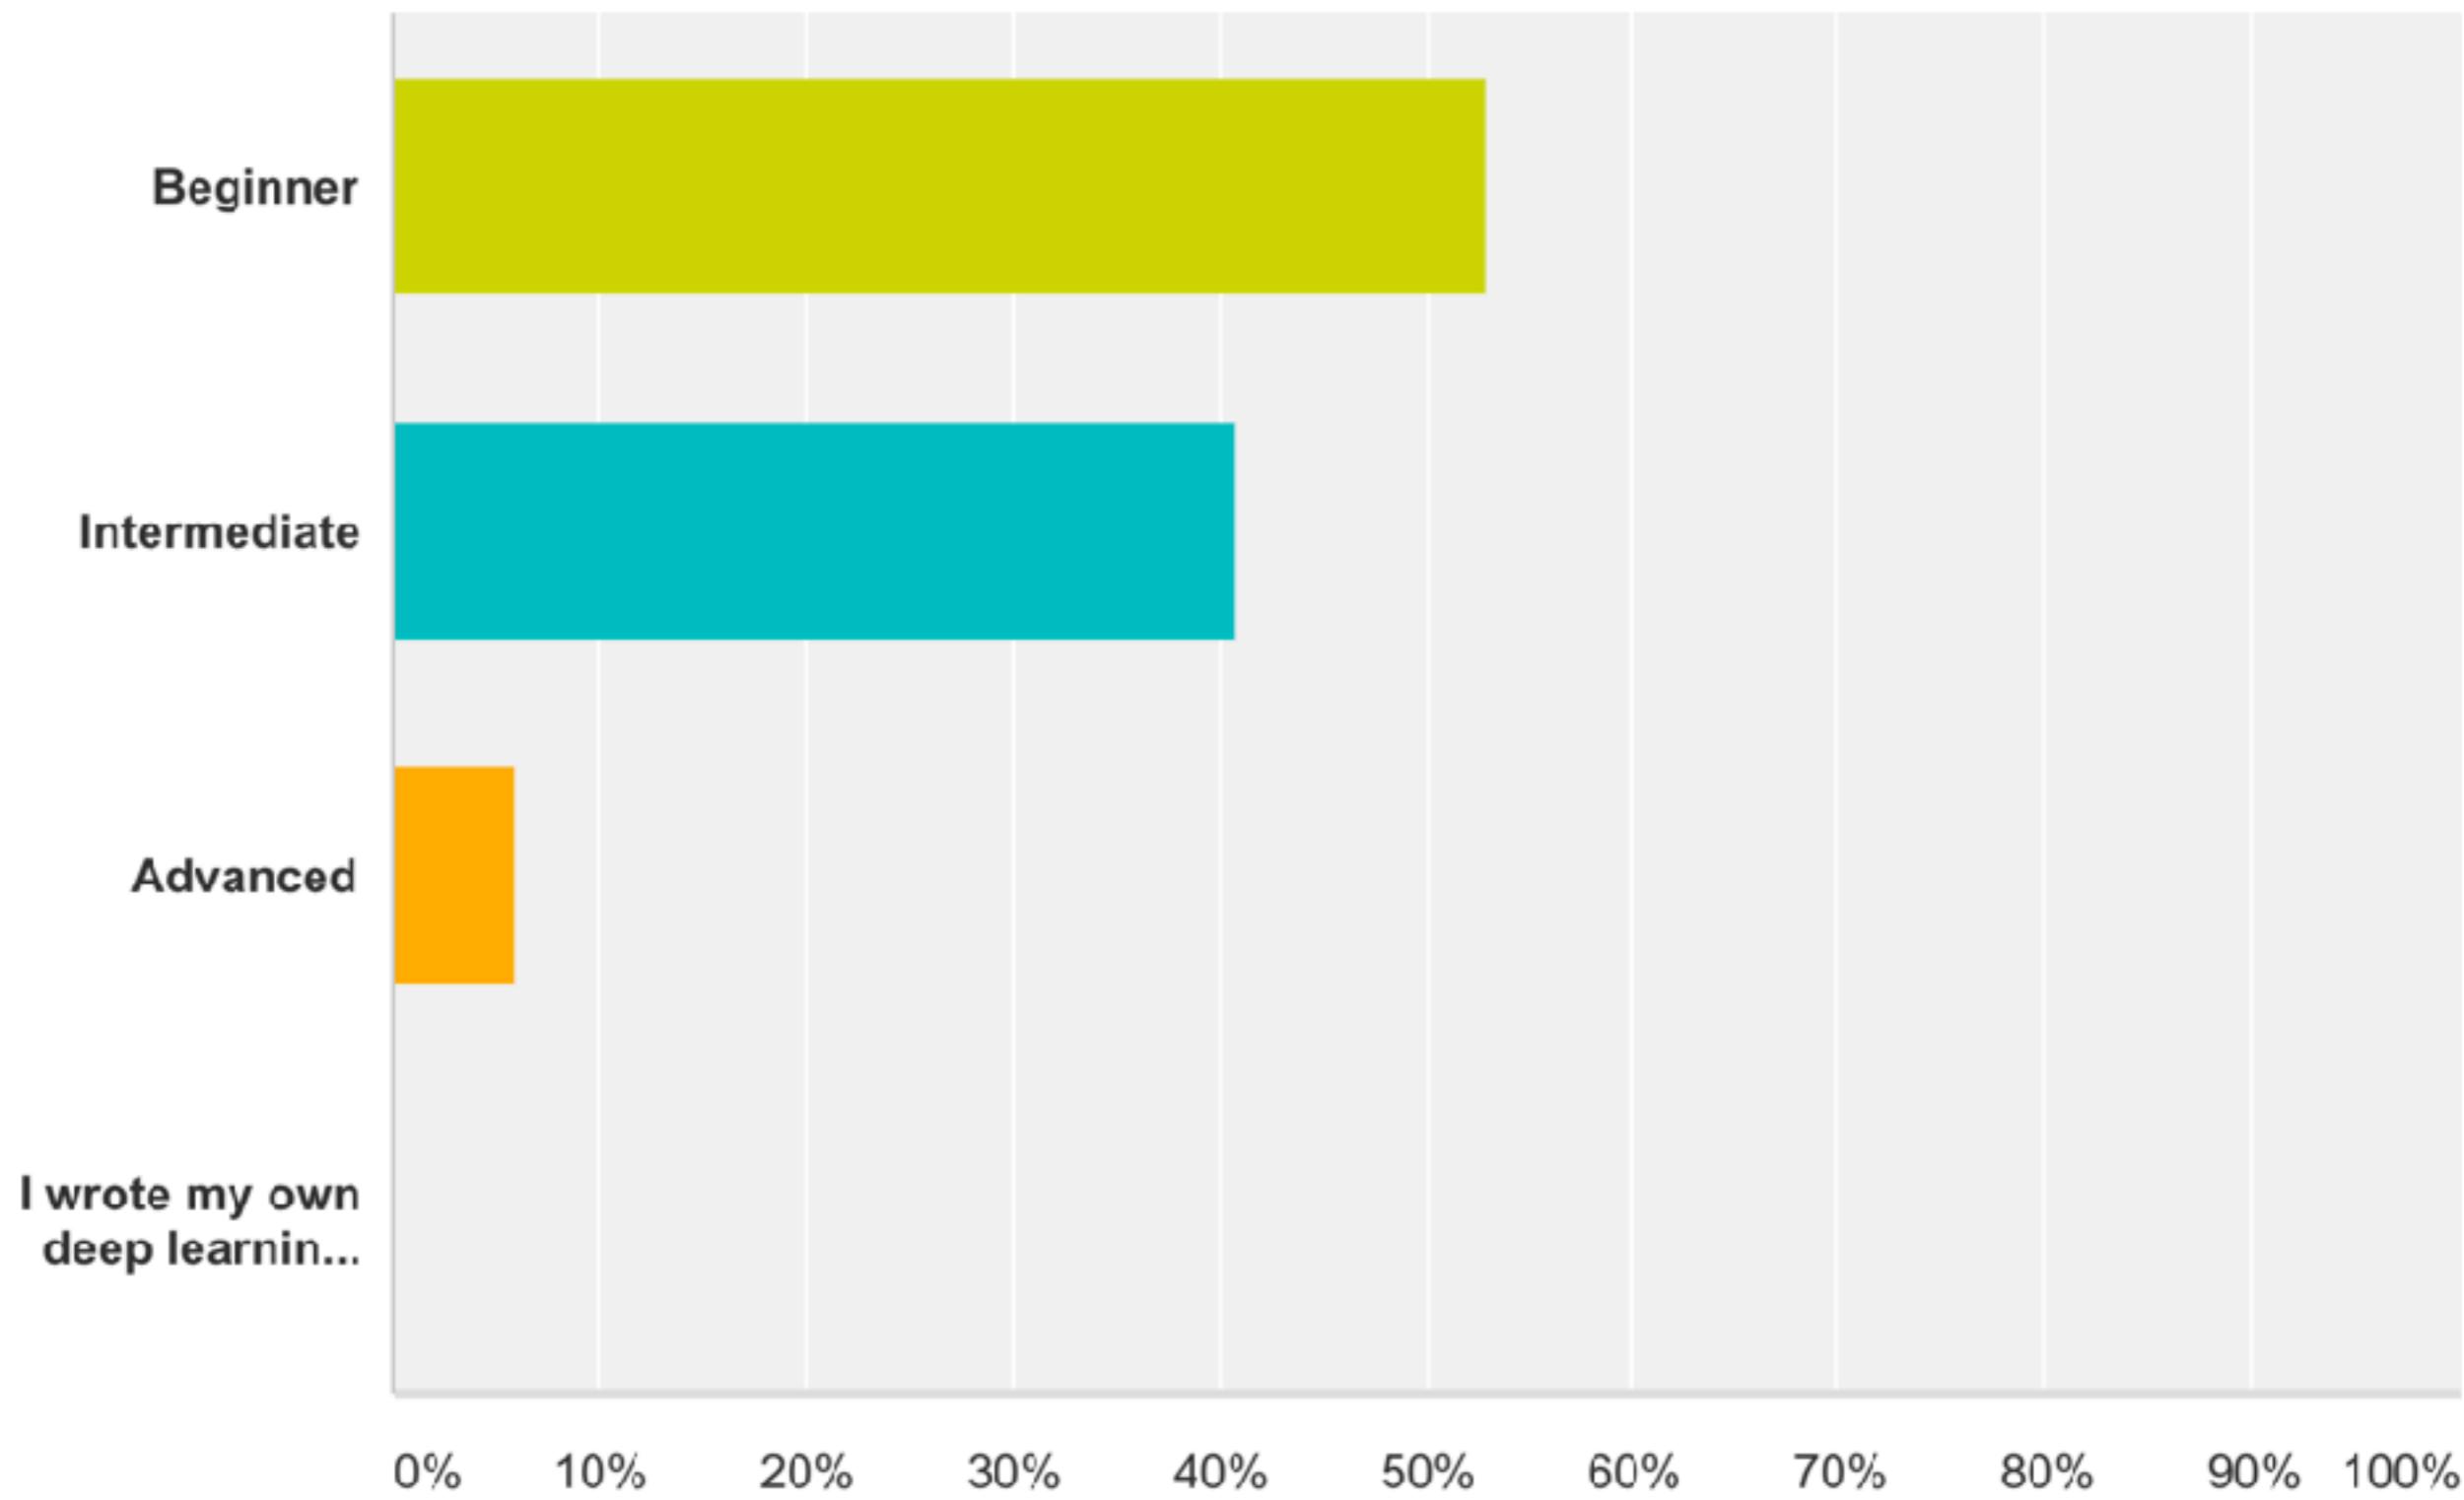
Which night of the week do you most prefer to meetup?

Answered: 47 Skipped: 2



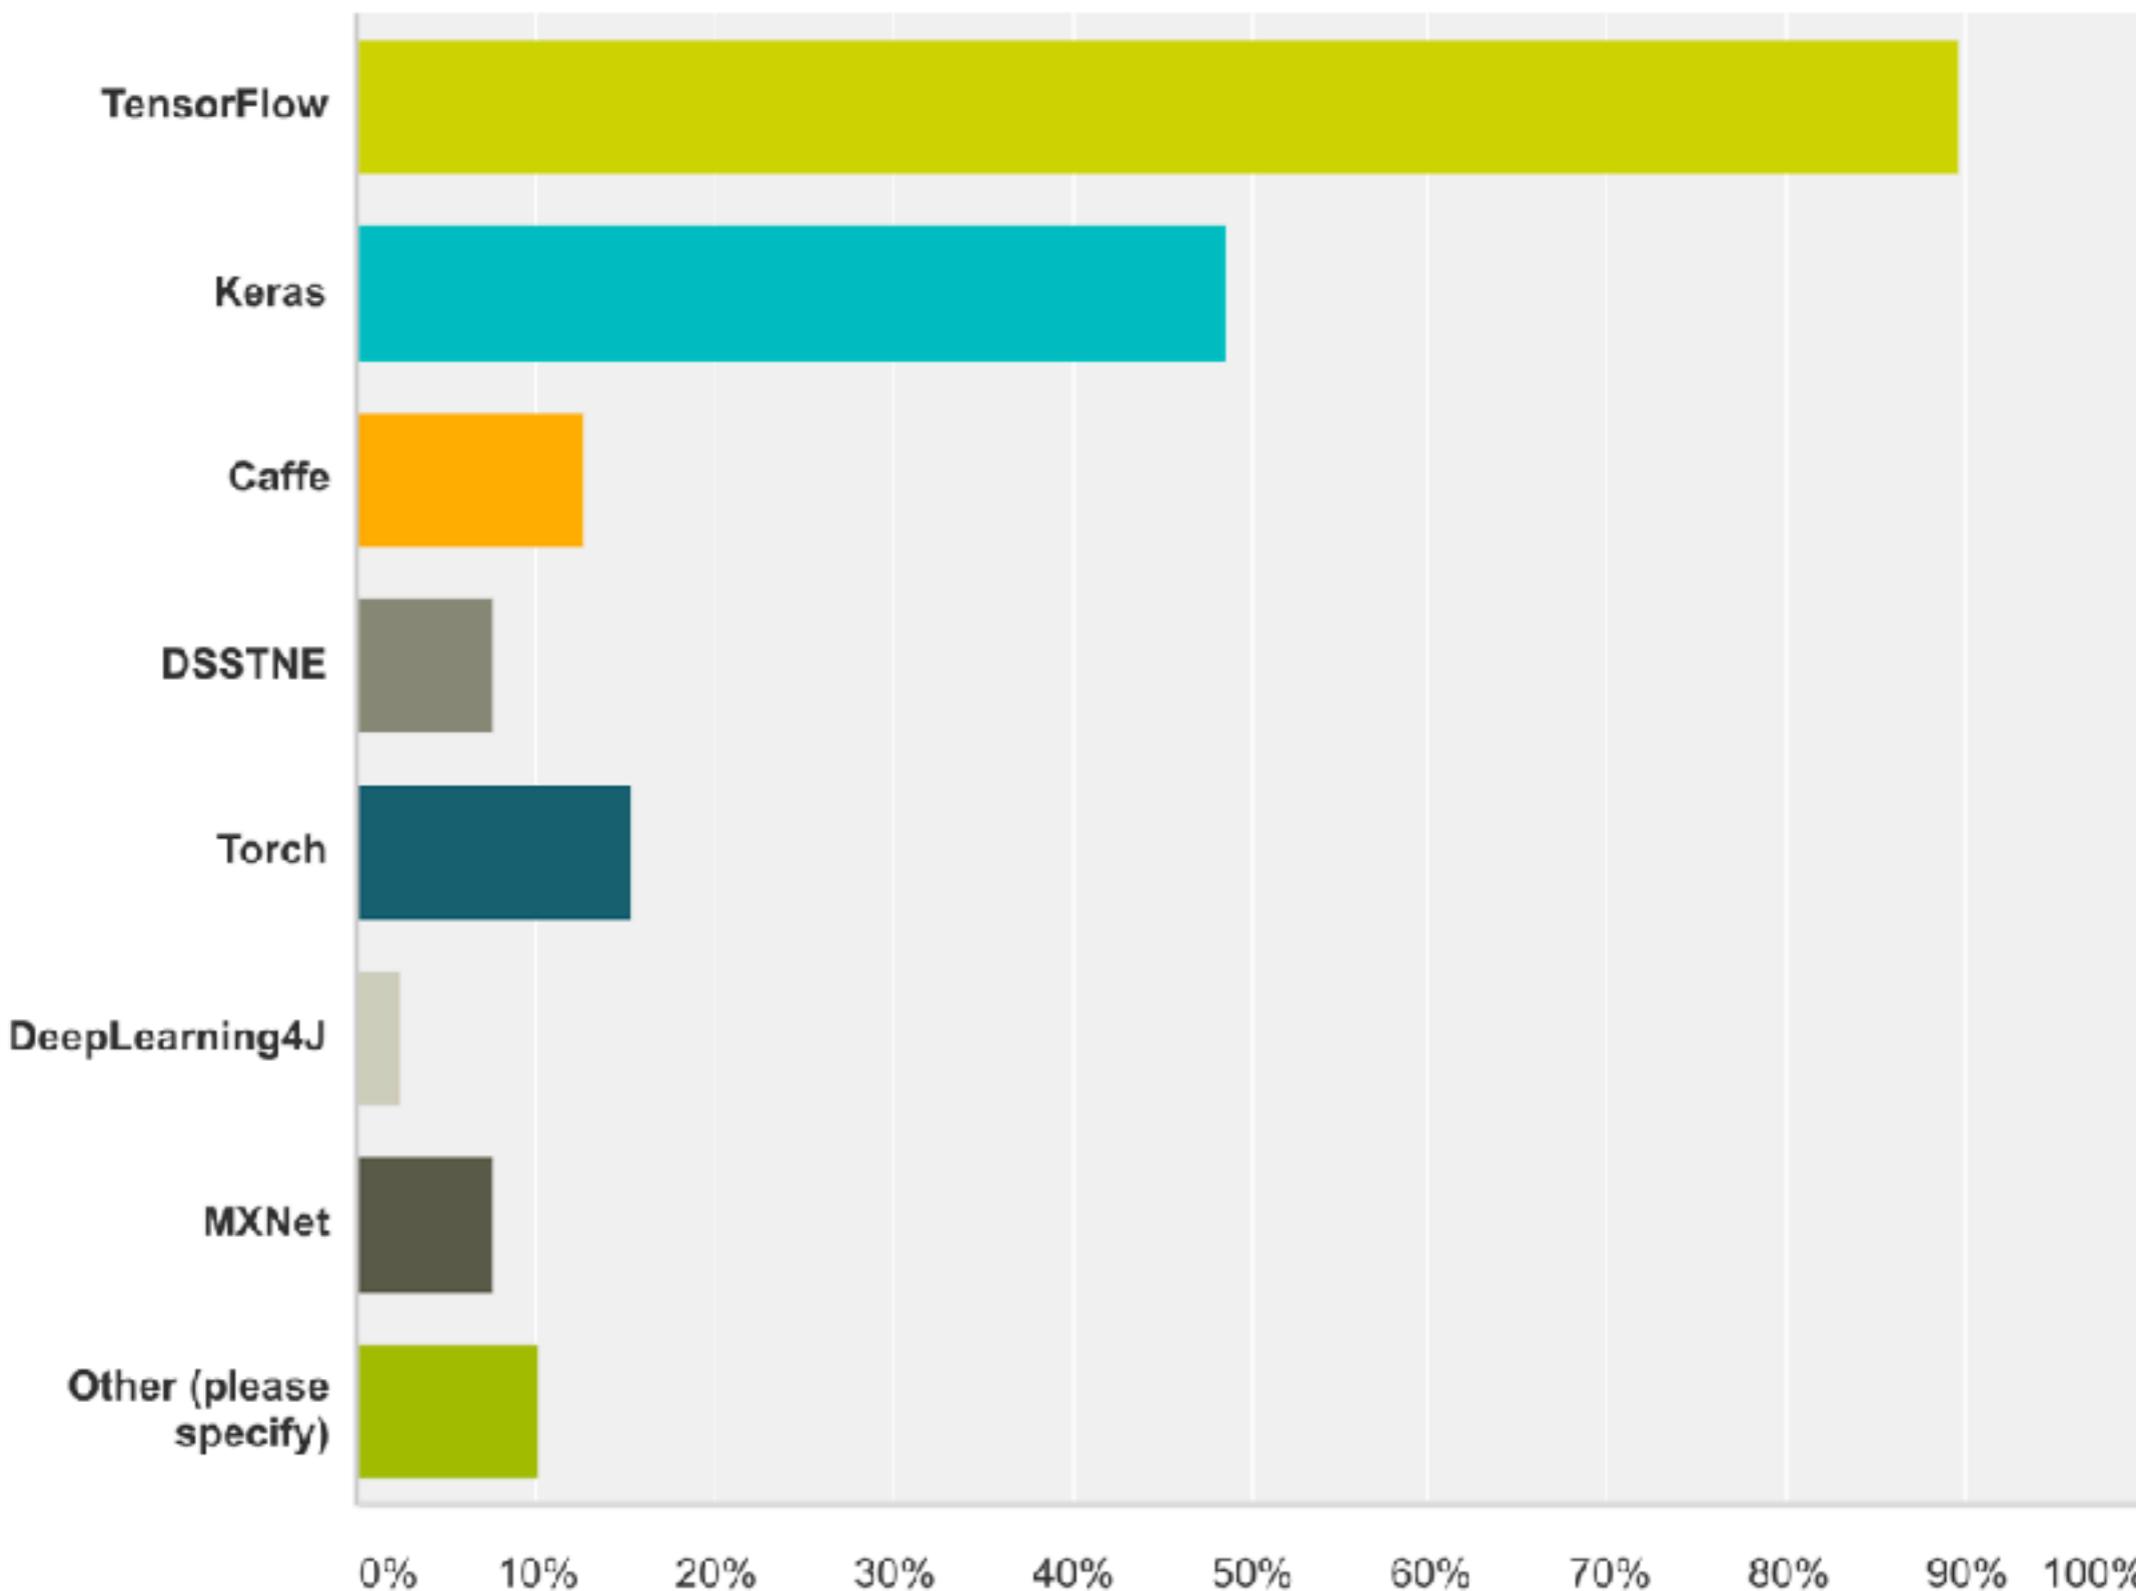
How would you rate your knowledge of deep learning?

Answered: 49 Skipped: 0



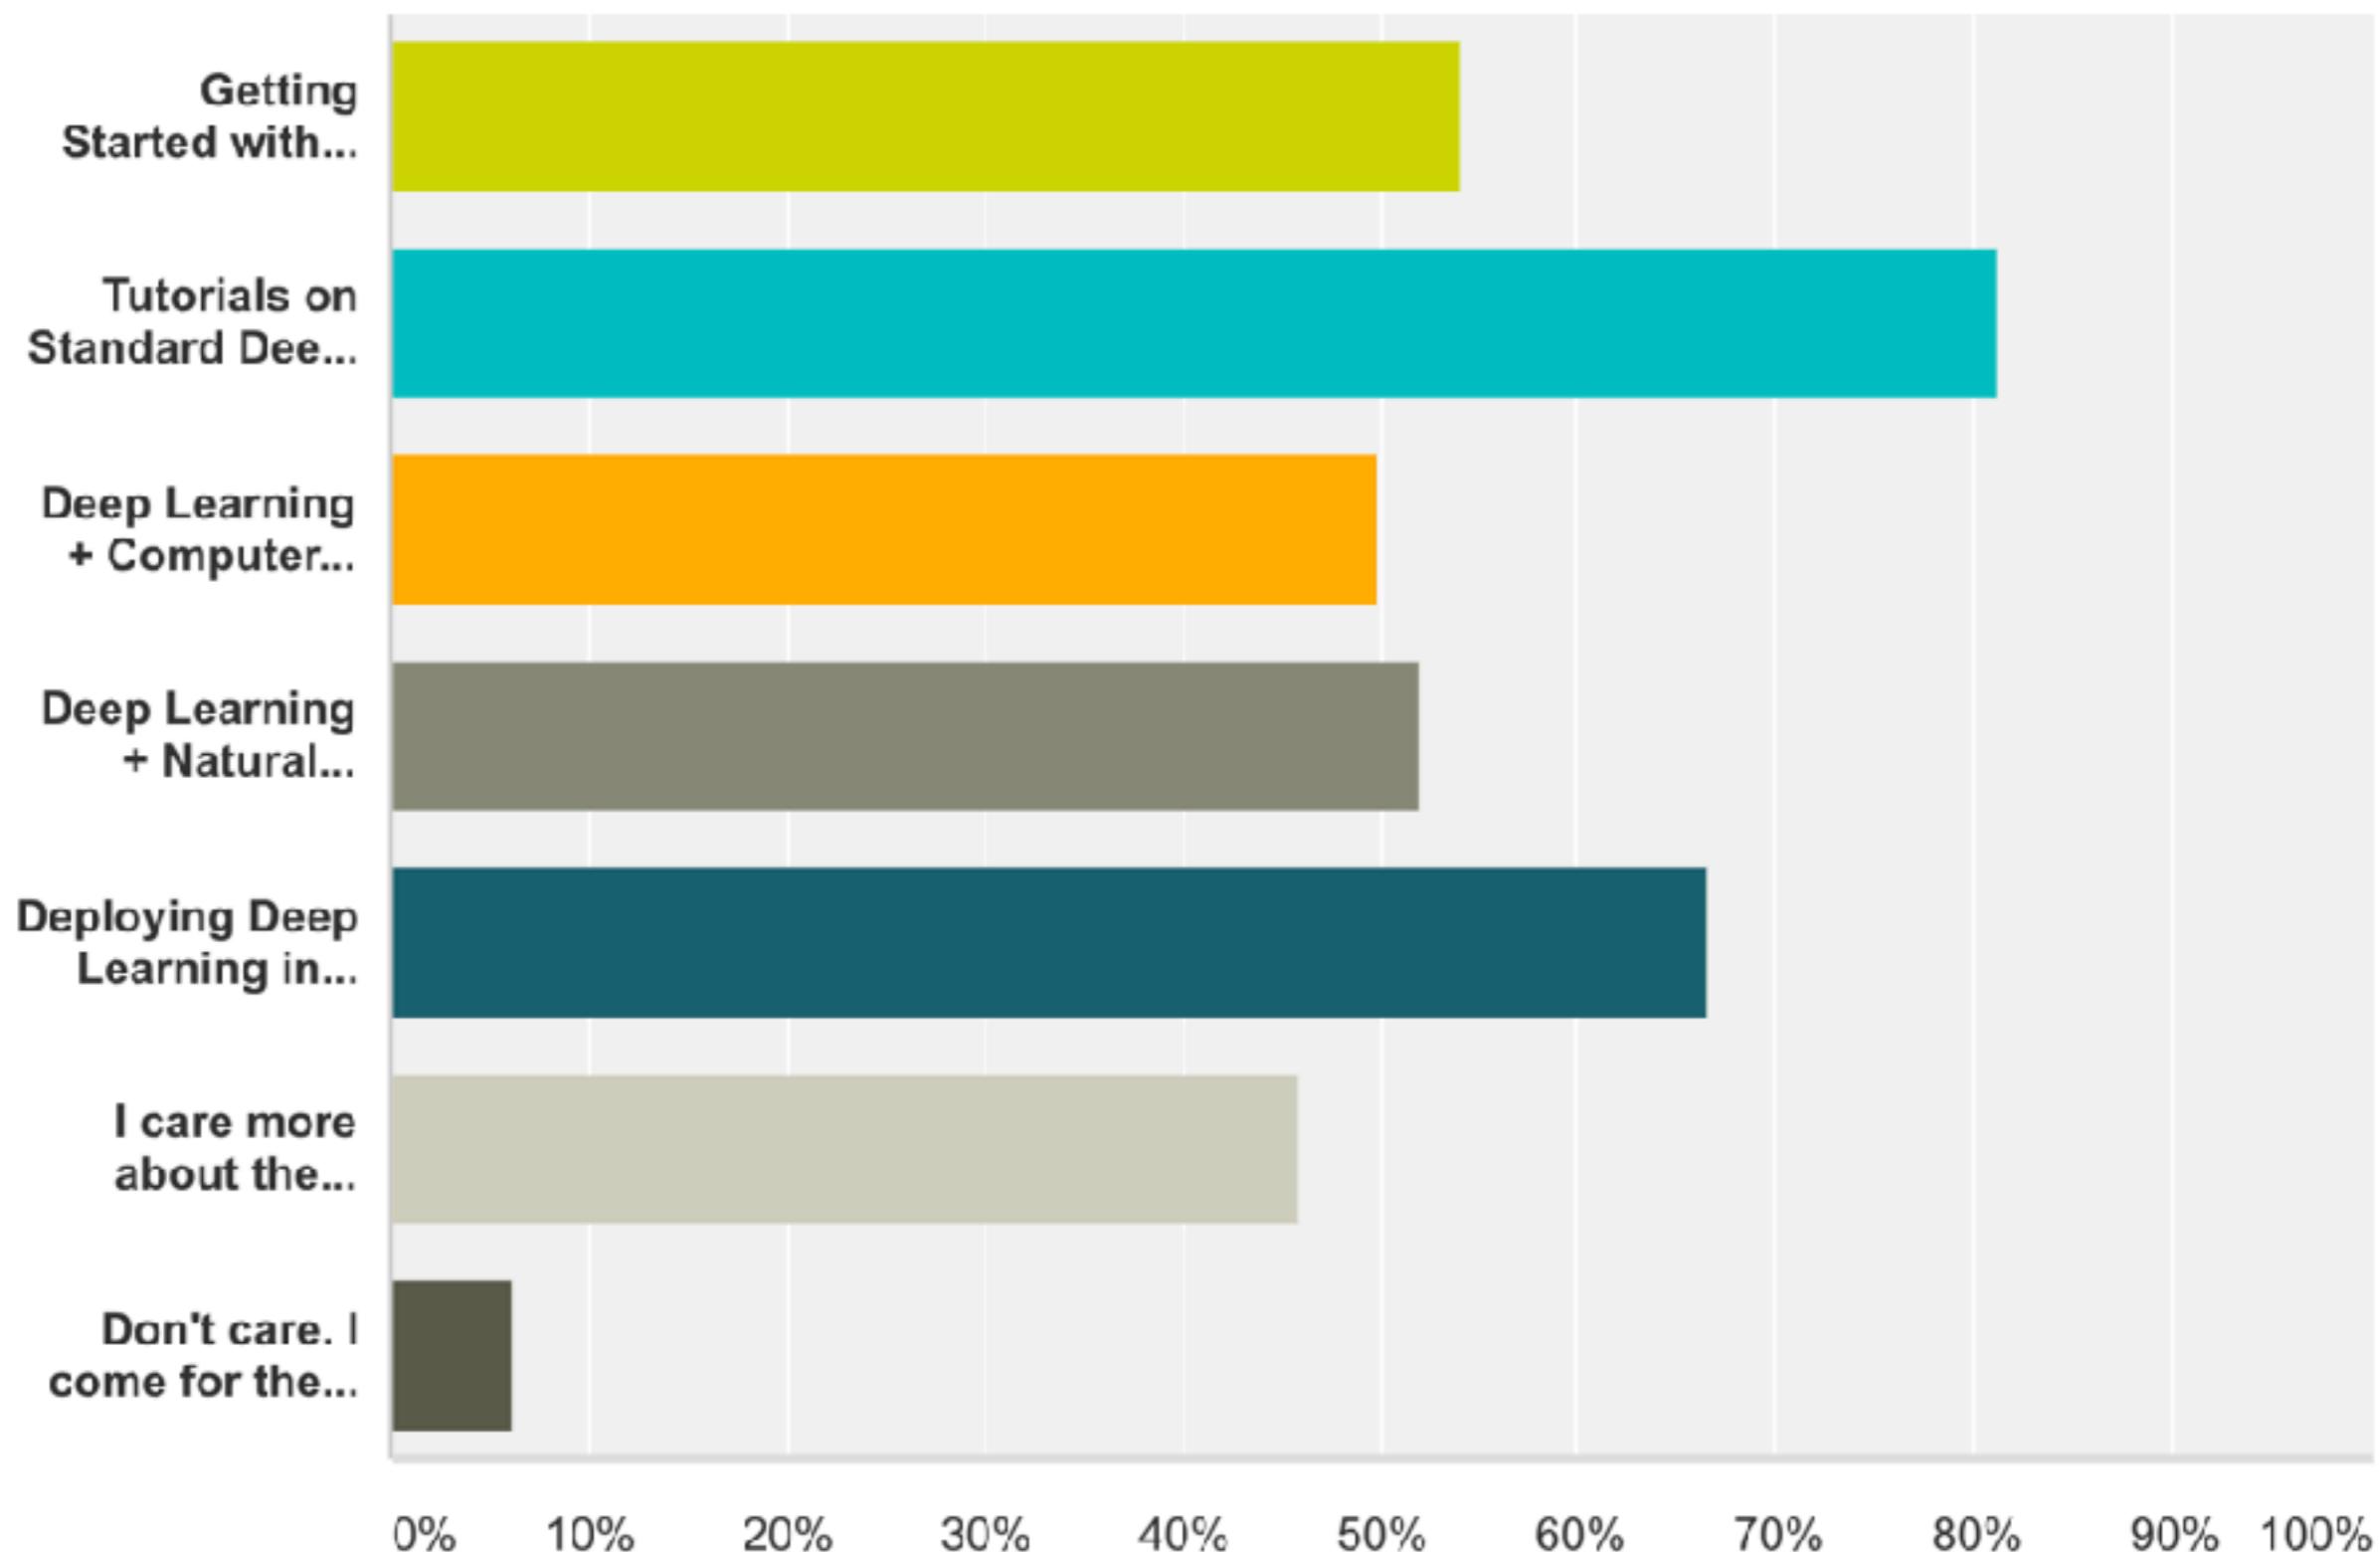
Which deep learning frameworks have you used?

Answered: 39 Skipped: 10



Which topics would you most prefer to hear?

Answered: 48 Skipped: 1



Answer Choices	Responses
▼ Getting Started with Deep Learning	54.17% 26
▼ Tutorials on Standard Deep Learning Networks (e.g., recurrent neural nets, convolutional neural nets)	81.25% 39
▼ Deep Learning + Computer Vision	50.00% 24
▼ Deep Learning + Natural Language Processing	52.08% 25
▼ Deploying Deep Learning in Production	66.67% 32
▼ I care more about the applications and what people are building.	45.83% 22
▼ Don't care. I come for the drinks and the fellowship.	6.25% 3

Total Respondents: 48

FEEDBACK / SUGGESTIONS

“

- * Practical Debugging of Deep Learning Networks
- * Tips and tricks to improve performance (such as vanishing gradients)
- * Advanced topics such as RNN with Hands-On and lecture

FEEDBACK / SUGGESTIONS

“

How to explain DL in more common terms and less math - concepts and processes tied to applications

FEEDBACK / SUGGESTIONS

“

More hands on meetup please.

Call for Talks

TOPICS

- Recurrent Neural Networks
- Image Caption Generation
- Semantic Segmentation
- Generative Adversarial Networks
- Deep Reinforcement Learning

OTHER IDEAS

- Blog posts
- Raspberry Pi
- Google “Awesome TensorFlow”





Slav Ivanov [Follow](#)
Entrepreneur / Hacker
May 29 · 16 min read

The \$1700 great Deep Learning box: Assembly, setup and benchmarks

Building a desktop after a decade of MacBook Airs and cloud servers

After years of using a thin client in the form of increasingly thinner MacBooks, I had gotten used to it. So when I got into Deep Learning (DL), I went straight for the brand new at the time Amazon P2 cloud servers. No upfront cost, the ability to train many models simultaneously and the general coolness of having a machine learning model out there slowly teaching itself.

However, as time passed, the AWS bills steadily grew larger, even as I switched to 10x cheaper Spot instances. Also, I didn't find myself training more than one model at a time. Instead, I'd go to lunch/workout/etc. while the model was training, and come back later with a clear head to check on it.

Hotdog!



Share

No Thanks

Not hotdog!



Share

No Thanks



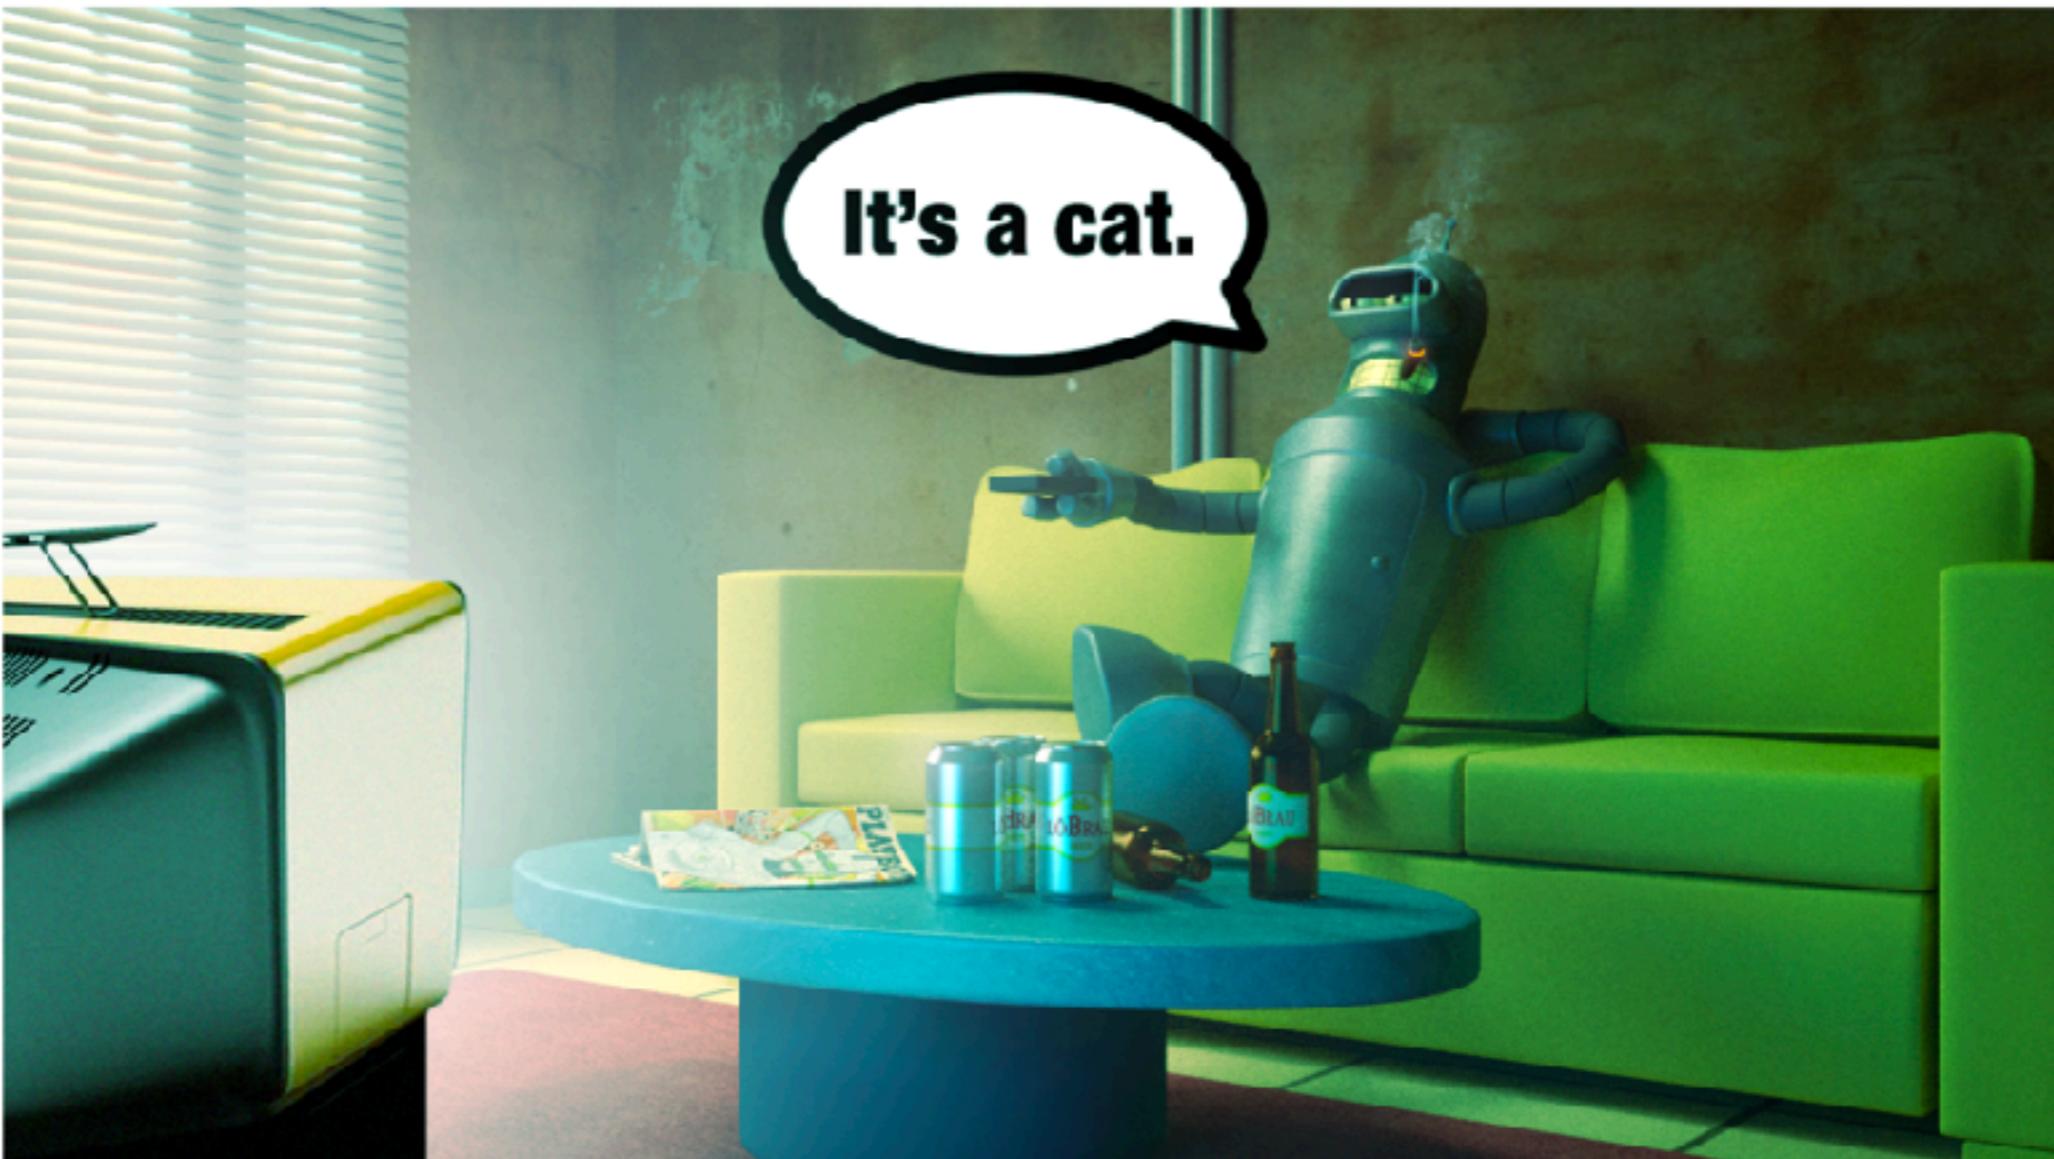
Matt Harvey [Follow](#)

Founder and CPO of Coastline Automation, where we're using AI to make every car autonomous and...

Mar 21 · 10 min read

Five video classification methods implemented in Keras and TensorFlow

Exploring the UCF101 video action dataset



[h/t [@joshumaule](#) and [@surlyrightclick](#) for the epic artwork.]